

Four-Month MVP Plan for an AI-Driven Pharmacy in Ontario

Executive summary

A four-month MVP is feasible **if it is framed as an AI-assisted pharmacy operations platform with pharmacist sign-off**, not as a fully autonomous public pharmacy. In Ontario, community pharmacies must be accredited by the Ontario College of Pharmacists, online pharmacy services must be operated by an accredited Ontario “bricks-and-mortar” pharmacy, and remote dispensing models still impose pharmacist-supervision requirements. For remote dispensing locations, Ontario permits automated pharmacy systems, but the technology must be approved by the College Board; if the pharmacist is not physically present at the remote site, there must be a live two-way audiovisual link, and the pharmacist must be physically present in the accredited pharmacy that operates the remote site. OCP also requires new community pharmacies and remote dispensing locations to submit accreditation materials at least 45 days before opening. ¹

That means the right MVP objective is **not** “open the world’s first fully AI-run pharmacy in 16 weeks.” The right objective is “**build and validate the core operating system for an AI-first pharmacy, inside a partner accredited Ontario pharmacy or in a closed mock dispensary lab, with one licensed Ontario pharmacist overseeing all clinically meaningful decisions.**” The MVP should handle prescription intake, authenticity triage, data extraction, structured pharmacist review, interaction screening, audit logging, basic inventory workflows, and a **demo-grade** dispensing hardware mockup for non-controlled oral solids only. Controlled substances, injectables, compounding, refrigeration-heavy workflows, insurance adjudication, and unsupervised patient-facing clinical advice should stay out of scope for the first 16 weeks. That narrowing is strongly supported by Ontario’s oversight model and by OCP’s specific restrictions on remote automated dispensing of narcotics and controlled drugs. ²

The strongest near-term thesis is this: start with **high-confidence operational automation plus pharmacist oversight**, not full autonomy. Build for the Ontario context by using accredited-pharmacy workflows, PHIPA-grade safeguards, electronic audit logging, and secure prescription transmission practices. PHIPA requires reasonable safeguards against theft, loss, unauthorized use, disclosure, copying, modification, and disposal of personal health information, and Ontario’s statute adds electronic audit-log obligations for custodians using electronic systems to handle PHI. PIPEDA still matters where commercial/private-sector processing or cross-border transfers arise, and HIPAA matters only if the product later serves U.S. covered entities or business associates. ³

A realistic 16-week output is a **pre-pilot MVP** with five validated capabilities:

- a secure prescription-ingestion pipeline;
- a prescriber-authentication and authenticity-triage workflow;
- a pharmacist dashboard with queueing, batch review where clinically appropriate, overrides, and audit logs;
- a drug normalization and interaction/contraindication engine backed by Canadian regulatory data plus a commercial interaction source;

and a basic dispensing demo that proves software-to-hardware orchestration for one narrow use case such as counting tablets from stock bottles into vials. 4

The biggest strategic adjustment, because of timing, is e-prescribing connectivity. Canada Health Infoway states that the **PrescribeIT operated service concludes on May 29, 2026**, while a **national ePrescribing standard** is now the path forward. So the MVP should not hard-wire itself to PrescribeIT as a product dependency. It should support: secure upload/fax ingestion immediately; registry-based prescriber verification; and an adapter layer for future connection to the national standard or jurisdiction/vendor-specific e-prescribing implementations. 5

Assumptions and regulatory frame

This report assumes the pilot is intended for **Ontario, Canada**, with Toronto as the founder’s base but not necessarily the first deployment site. Because budget, exact pilot site, partner pharmacy access, and pharmacist availability were not specified, the plan assumes the following: an accredited Ontario pharmacy partner or a near-term signed pilot LOI; one licensed Ontario pharmacist available for design review and weekly UAT; no existing codebase; and a lean startup build budget in the rough range of **CAD 200,000 to CAD 350,000** for software, contractors, hardware prototypes, legal/privacy review, and cloud, excluding actual public pharmacy fit-out or full accreditation build-out. This budget figure is a planning assumption, not a sourced market benchmark.

For Ontario, several regulatory facts directly shape the MVP. OCP regulates both registrants and pharmacies, and pharmacy operations must meet legislation, standards, and accreditation requirements. Online pharmacy services in Ontario must be tied to an accredited pharmacy, and OCP’s internet-site policy requires authenticity controls for prescription orders, compliant transmission methods, integrated recordkeeping, and homepage disclosure of accreditation details, designated manager, address, phone number, and pharmacist consultation hours. Health Canada likewise warns that legitimate online pharmacies should require a valid prescription and have a Canadian-licensed pharmacist available. 6

Prescription authenticity deserves special emphasis. OCP states that ordinary unsecured email is **not** a secure medium for transmitting prescriptions, and that when a prescription is emailed insecurely, pharmacists are encouraged to contact the prescriber to verify authenticity and move to established channels such as fax, verbal confirmation, or secure electronic transmission. That means the MVP’s authenticity workflow should treat plain emailed or patient-forwarded prescriptions as **unverified intake**, not as automatically valid electronic prescriptions. 7

The cleanest MVP scope therefore is:

Scope decision	Recommendation	Why it belongs in the MVP boundary
Pilot setting	Closed pilot inside partner Ontario pharmacy or controlled mock dispensary	Aligned with Ontario accreditation and supervision realities; avoids premature public operation claims. 8

Scope decision	Recommendation	Why it belongs in the MVP boundary
Supervising clinician	One licensed Ontario pharmacist, onsite in the parent pharmacy for any remote/automated pilot behavior	Consistent with Ontario remote-dispensing rules. ⁹
Medication scope	Non-controlled oral solid medications in stock bottles only	Safer, simpler hardware integration; avoids controlled-drug/RDL restrictions and sterile workflows. ¹⁰
Prescription sources	Paper, faxed, scanned PDF/image, secure upload, and standard-ready e-prescribing adapter	Practical post-PrescribeIT path. ⁵
Out of scope	Narcotics/controlled substances, compounding, injectables, refrigeration tracking, public autonomous dispensing, insurance adjudication	Too much regulatory and operational risk for 16 weeks. ¹¹

One more boundary matters for product claims. Doctor handwriting and free-text notes remain a difficult technical problem. Current research continues to describe handwritten prescriptions as hard because of irregular spacing, handwriting variability, and non-standard abbreviations; healthcare document-understanding work also shows that hybrid systems combining learned extraction with domain-specific rules outperform pure model pipelines in real operational settings. In practice, that means the MVP should automate extraction and triage, but the pharmacist must remain the release authority whenever handwriting confidence is low or clinical ambiguity exists. ¹²

MVP scope and system architecture

The MVP should be built as an **operations OS** sitting between prescription intake, pharmacist review, inventory, and a narrow hardware demo. Health Canada's Drug Product Database is the basic Canadian product source of truth: it is maintained by Health Canada, updated nightly, and includes authorized drugs, DINs, and product monographs. It is a strong foundation for product normalization, DIN validation, and label/product lookups, but it is **not enough by itself** for robust drug-drug interaction screening, so a commercial clinical content source is still needed. ¹³

Prioritized feature list

Priority	Module	What the MVP should do
Must-have	Prescription intake and OCR	Ingest images/PDFs/faxes/uploads; classify source; extract text; assign confidence; preserve source image; create a reviewable structured draft.

Priority	Module	What the MVP should do
Must-have	Prescriber authentication	Verify physician and NP identity against Ontario registries; store verification evidence; quarantine unverified prescriptions; support callback workflow.
Must-have	Prescription parsing	Extract patient, prescriber, medication, strength, SIG, quantity, repeats, date, notes, and transmission channel; surface ambiguity flags.
Must-have	Drug interaction and contraindication engine	Normalize products/ingredients and generate structured alerts, severity, rationale, and management notes for pharmacist review.
Must-have	Pharmacist dashboard	Queue, batch triage, case detail view, override/reject flow, activity/audit logs, turnaround metrics, exception inbox, and signed approval record.
Must-have	Security and privacy	MFA, RBAC, encryption, secure object storage, immutable audit logs, PHI-safe logging, breach workflow, retention controls.
Must-have	QA and regulatory validation	Gold dataset, extraction benchmarks, negative/fraud tests, UAT scripts with pharmacist, SOP drafts, incident-reporting workflow.
Should-have	Inventory and auto-ordering	Basic stock decrement, reorder thresholds, low-stock alerts, wholesaler CSV/export or API integration placeholder.
Should-have	Patient chatbot and counseling	FAQ, refill readiness, approved counseling snippets, monograph-backed answers, and mandatory pharmacist escalation for anything clinical or uncertain.
Should-have	Basic dispensing hardware mockup and integration	Trigger one narrow tablet-dispense demo from approved script to device controller; confirm count result and pack event in software.
Nice-to-have	Analytics and staffing insights	Queue time, exception rate, handwriting-confidence trend, stockout forecast.
Nice-to-have	Voice/video counseling	Integrated call handoff to pharmacist and required consent flows for remote counseling.

The “must-have” set is driven by Ontario’s need for prescription authenticity controls, accredited-pharmacy operation, privacy safeguards, and auditable professional oversight. The “should-have” set is where the AI-pharmacy story starts to look differentiated to investors and pilot partners without creating avoidable first-version risk. ¹⁴

Recommended architecture

The most practical architecture is a **document-first, human-in-the-loop workflow**:

1. **Ingestion service** receives PDF/image/fax upload and stores the original artifact in encrypted object storage.
2. **OCR/document analysis** extracts text, boxes, and tables and returns per-field confidence.
3. **Document understanding layer** applies a hybrid parser: rules plus a document model to map OCR output into structured prescription data.
4. **Authentication service** verifies the prescriber against Ontario registries and checks transmission policy.
5. **Medication normalization** maps extracted drug names to DIN/product concepts using Health Canada DPD, with secondary normalization support where needed.
6. **Clinical screening** calls the interaction/contraindication engine and produces alerts for pharmacist review.
7. **Pharmacist workbench** shows the case, source image, extracted fields, auth status, alert stack, and approve/reject/escalate actions.
8. **Inventory/dispense orchestration** decrements stock and, for demo-eligible items only, may send a counted-dispense job to the hardware controller.
9. **Audit and incident layer** logs every view, change, approval, override, and dispense event.

This design is aligned with PHIPA's safeguards and electronic-audit-log expectations, and it is robust against one of the biggest operational truths in pharmacy automation: **ambiguity is expected**, not exceptional.

15

Team, hiring plan, stack, and vendor choices

Required team composition

A realistic lean team for 16 weeks is:

Role	FTE or contractor shape	Main responsibility
Founder / product lead	0.75–1.0 FTE	Scope control, pilot partner, workflows, weekly program management
Licensed Ontario pharmacist / clinical lead	0.3–0.5 FTE contractor	Clinical rules, SOPs, dataset labeling, UAT, approval workflow design
Tech lead / backend engineer	1.0 FTE	Core architecture, APIs, integrations, security foundations
ML / document AI engineer	1.0 FTE	OCR pipeline, parsing, extraction validation, confidence scoring
Full-stack engineer	1.0 FTE	Dashboard, intake UI, auth flows, ops tooling

Role	FTE or contractor shape	Main responsibility
Robotics / embedded engineer	0.5–0.75 FTE	Device controller, sensors, dispense mockup, hardware integration
Security / privacy lead	0.2–0.3 FTE fractional consultant	PHIPA/PIPEDA design review, threat model, policies, logging, incident response
QA / validation engineer	0.5 FTE from week 6 onward	Test plans, regression pack, gold-set evaluation, release readiness
UX designer	0.2–0.3 FTE during weeks 1–6	Workflow usability for pharmacists and ops staff

If capital is tight, the first hires should be: **tech lead, ML/document AI engineer, pharmacist contractor, and one full-stack engineer**. Robotics can be contracted after week 5 once the data flow from approved prescription to dispense job is clear.

Hiring and contractor plan

Use employees or long-term contractors for product, core engineering, and ML. Use fractional experts for privacy/security and pharmacy jurisprudence. The pharmacist contractor should be brought in immediately, not after build, because Ontario workflow design, labeling assumptions, exception routing, and counseling boundaries are too domain-specific to retrofit cheaply. OCP's standards, remote supervision model, and medication-safety expectations make early pharmacist involvement a design requirement, not a compliance afterthought. ¹⁶

Recommended tech stack

A strong default stack for a zero-codebase startup MVP is:

Layer	Recommendation	Why
Frontend	Next.js + TypeScript	Fast product velocity, good admin UI ergonomics
Backend APIs	Python + FastAPI	Best fit for OCR, parsing, clinical-rule services, and rapid API development
Async jobs	Celery or Temporal	Reliable long-running document workflows and retries
Database	PostgreSQL	Strong transactional core for orders, queues, audit refs
Cache / queues	Redis	Simple high-speed queue/cache layer
Object storage	S3-compatible or Azure Blob in Canadian region	Encrypted PHI artifact storage
Search / retrieval	PostgreSQL full text first; OpenSearch later if needed	Avoid overbuilding

Layer	Recommendation	Why
Auth	Managed enterprise IdP with MFA and SSO; enforce RBAC	Security and admin control
Audit log store	Append-only event log plus warehouse export	Supports PHIPA auditability
Cloud region	Canadian region by default	Operationally helpful for residency and procurement discussions; AWS offers Canada (Central), Azure offers Canada Central/Canada East, and Google Cloud offers Toronto and Montréal regions. ¹⁷

For document understanding, the best practical path is **managed OCR plus a hybrid parser**. Research on LayoutLMv3, Donut, and TrOCR shows strong document-understanding and text-recognition capability, but in healthcare operations, hybrid pipelines that combine document models with rules are often more reliable than end-to-end “magic” systems. The right early build is therefore: managed OCR for robustness, rule-based extraction for obvious patterns, model-assisted extraction for messy regions, and mandatory pharmacist review when confidence drops below threshold. ¹⁸

Data sources and prioritized APIs

Priority	Source / API	Recommended use
Highest	Health Canada Drug Product Database and API/data extracts	DIN/product lookup, product normalization, product monographs, authorized-product validation. Updated nightly. ¹⁹
Highest	CPSO Physician Register	Verify Ontario physicians and CPSO numbers. ²⁰
Highest	CNO Find a Nurse / Nurse Renewal Check	Verify Ontario nurse practitioners and support bulk-verification patterns. ²¹
High	DrugBank, FDB, or Medi-Span	Drug-drug interactions, screening content, severity, management recommendations. ²²
Medium	RCDSO Find a Dentist	Add only when dental prescriptions enter MVP scope. ²³
Medium	Infoway national ePrescribing standard	Future interoperability layer after PrescribeIT service conclusion. ²⁴
Medium	RxNorm / RxClass	Supplemental normalization and classes; not sufficient for interaction screening because RxNav drug-interaction features were discontinued. ²⁵

OCR vendor options

Option	Public cost signal	Pros	Cons
Amazon Textract	Detect text example pricing shows \$0.0015/page for first 1M pages; forms \$0.05/page ; tables \$0.015/page in AWS examples. ²⁶	Strong OCR plus form/table extraction; simple developer workflow; fast startup execution	Handwriting and clinician free text still require downstream validation; cost can climb on form-heavy workloads
Google Document AI	OCR processor \$1.50/1,000 pages ; Form Parser \$30/1,000 pages ; Custom Extractor \$30/1,000 pages . ²⁷	Clear pricing; strong document processing ecosystem; good custom parsers	Can become expensive if you rely heavily on form parsing/custom extraction
Azure AI Document Intelligence	Public pricing page states documents are billed by pages analyzed; free tier exists; commitment tiers include prebuilt \$8,640/year for 1.2M pages and custom extraction \$24,480/year for 1.2M pages for disconnected containers. ²⁸	Enterprise-friendly, strong Microsoft ecosystem and identity integration	Public pricing is less transparent in some views; real costs often require calculator or quote
Tesseract	Open-source, no license fee. ²⁹	Free, self-hostable, useful baseline and fallback	Weak choice for messy handwritten prescriptions without substantial tuning; higher ops burden

Recommendation: use **managed OCR first** for the MVP. Tesseract is fine as a benchmark or low-cost fallback, but it should not be the primary engine for safety-critical pharmacy extraction in a 16-week launch.

Drug interaction vendor options

Option	Cost	Pros	Cons
RxNav / RxNorm	Free for normalization APIs, but RxNav's interaction features were discontinued in 2024. ³⁰	Excellent for normalization and class lookups; official U.S. NLM resource	Not suitable as the MVP's interaction engine
DrugBank Clinical API	Enterprise quote-based; claims 1.3+ million drug-drug interactions with severity, evidence, and management fields. ³¹	Rich interaction content; developer-friendly API surface; fast startup integration	Commercial licensing cost; region/licensing negotiation needed

Option	Cost	Pros	Cons
FDB MedKnowledge / Drug-Drug Interaction Module	Enterprise quote-based. ³²	Widely used in healthcare IT; configurable severity and alert filtering; U.S. and Canada coverage	Quote-led sales cycle; integration work can be heavier
Wolters Kluwer Medi-Span Clinical APIs	Enterprise quote-based. ³³	Strong clinical screening capabilities, including patient- specific alerting concepts	Enterprise procurement cycle; more complex implementation and optimization

Recommendation: for a Canadian startup MVP, **DrugBank** or **FDB** are the practical shortlist. Use Health Canada DPD as the product source of truth and a commercial DDI source for clinical screening.

Robotics prototyping options

Option	Public cost signal	Pros	Cons
Elephant Robotics myCobot 280	\$599–\$849 depending on compute variant; suction pump about \$149.99 ; adaptive gripper about \$129.99 . ³⁴	Cheap, accessible, excellent for founder demo and software- hardware proof of concept	Light payload, lower precision and durability; not near production pharmacy automation
UFACTORY Lite 6	About \$3,389.92 without gripper; vacuum gripper about \$302.40 . ³⁵	Better reach, sturdier integration platform, useful for repeatable tabletop pick-and-place	Higher cost; still a prototype desk robot, not validated dispensing equipment
DOBOT Magician E6	\$6,995 including suction gripper on official U.S. product page. ³⁶	Desktop cobot with stronger industrial feel; better partner-demo optics	Most expensive of the three; still not pharmacy- certified dispensing technology


Recommendation: start with **myCobot 280** if budget discipline matters, or **Lite 6** if the goal is a sharper enterprise demo. In either case, position the hardware as a **mock dispensing demonstrator**, not as production-safe pharmacy automation.

Sixteen-week project plan

Role legend

For the weekly plan below: **PL** = product lead, **CL** = pharmacist/clinical lead, **BE** = backend engineer, **MLE** = ML engineer, **FE** = full-stack engineer, **RE** = robotics engineer, **SEC** = security/privacy lead, **QA** = QA engineer.

Gantt-style view



Weekly milestones

Week	Key deliverables	Owner roles	Est. effort	Dependencies	Success criteria
Week 1	Finalize scope, pilot assumptions, partner workflow map, medication scope, target prescription formats, architecture draft	PL, CL, BE, SEC	140 hrs	None	Signed scope memo; out-of-scope list approved; 10–15 representative prescription samples collected
Week 2	Threat model, PHIPA/PIPEDA control list, data model, ingestion API skeleton, UI wireframes	BE, SEC, FE, CL	160 hrs	Week 1	Security design review complete; source-image retention and audit event schema agreed
Week 3	OCR spike with 2–3 engines; benchmark typed, faxed, and messy samples; select OCR vendor	MLE, BE, CL, QA	180 hrs	Weeks 1–2	OCR decision made; baseline extraction benchmark documented
Week 4	Registry-auth design; CPSO/CNO verification workflow; fraud/quarantine logic; go/no-go gate	BE, FE, CL, SEC	190 hrs	Weeks 1–3	Gate A: partner pharmacist committed, registry workflow proven, dataset access confirmed
Week 5	End-to-end intake pipeline from upload to structured draft; confidence scoring; manual correction UX	MLE, BE, FE	210 hrs	Weeks 3–4	User can upload prescription and see structured draft with field confidence and source image

Week	Key deliverables	Owner roles	Est. effort	Dependencies	Success criteria
Week 6	Clinical schema finalized; Health Canada DPD integration for DIN and product lookup; QA gold set begins	BE, MLE, CL, QA	220 hrs	Week 5	≥90% typed-field recall on gold set for patient/prescriber/drug core fields
Week 7	Commercial DDI engine integrated; alert rendering in dashboard; override reasons and escalation states	BE, FE, CL	220 hrs	Week 6	Drug alerts visible with severity/rationale; pharmacist can accept, override, or escalate
Week 8	Pharmacist dashboard v1; queue, case detail, structured review, signed approval; go/no-go gate	FE, BE, CL, QA	220 hrs	Weeks 5–7	Gate B: end-to-end review flow works; typed/faxed scripts hit target accuracy; auth failures are quarantined
Week 9	Inventory v1; stock decrement, reorder thresholds, low-stock alerts, dispense eligibility rules	BE, FE, CL	200 hrs	Week 8	Approved scripts update stock and create dispense task only for allowed SKUs
Week 10	Counseling module v1 using approved scripts and monograph-backed retrieval; pharmacist escalation button	FE, MLE, CL	220 hrs	Weeks 6–8	Chatbot answers only from approved content and reliably escalates uncertain clinical questions
Week 11	Hardware BOM, controller code, dispense simulation or bench rig, command protocol from backend	RE, BE	220 hrs	Weeks 8–9	Software sends approved task to mock hardware controller and receives completion/error event
Week 12	Hardware demo v1; audit-log viewer; incident logging workflow; go/no-go gate	RE, BE, QA, SEC, CL	220 hrs	Weeks 10–11	Gate C: approved script can trigger a narrow dispense demo; audit events complete and queryable
Week 13	Security hardening, MFA, RBAC, secrets, log scrubbing, backup/restore, breach-response drill	SEC, BE, FE	180 hrs	Weeks 2–12	Security checklist closed; no PHI in app/debug logs; backup restore tested

Week	Key deliverables	Owner roles	Est. effort	Dependencies	Success criteria
Week 14	UAT with pharmacist on full scenarios; edge-case fixes; SOP drafts; training material for pilot users	CL, QA, FE, BE	180 hrs	Weeks 8–13	Pharmacist completes test script with acceptable friction; SOP pack drafted
Week 15	Dry-run pilot in mock dispensary or partner site; reliability testing; metrics dashboard	PL, QA, RE, BE, FE, CL	160 hrs	Weeks 11–14	20–50 scripted test cases run end to end; error backlog triaged by severity
Week 16	Demo day build, KPI report, risk review, pilot readiness decision, next-phase plan	PL, CL, BE, QA, SEC	180 hrs	All prior weeks	Gate D: ship/no-ship decision based on readiness metrics and risk thresholds

A realistic target for the 16-week build is about **3,300 team-hours**. That is lean but credible if scope discipline is maintained and the hardware stays narrow.

Risk register and decision gates

Risk register

Risk	Category	Likelihood	Impact	Mitigation	Trigger / watchpoint
Handwritten prescriptions are too noisy for safe extraction	Technical / clinical	High	High	Restrict MVP to typed, faxed, or high-quality scans for primary path; use pharmacist review for all low-confidence fields; collect partner-specific training data	Core-field precision on handwritten set stays below threshold by week 8
Prescriber-auth workflow is brittle or slow	Technical / regulatory	Medium	High	Start with CPSO + CNO only; support manual callback verification; log evidence; defer full multi-college coverage	Too many prescriptions land in unverified state or manual callback queue

Risk	Category	Likelihood	Impact	Mitigation	Trigger / watchpoint
Commercial DDI vendor procurement slips	Technical / procurement	Medium	High	Run vendor selection in weeks 1–3; keep DPD normalization independent; use placeholder alert framework until license live	Contract not signed by week 6
Privacy/ security design is bolted on too late	Security / regulatory	Medium	High	Fractional privacy lead from week 1; threat model in week 2; PHI-safe logging and MFA before beta	PHI appears in logs or backups during testing
Product overreaches into public autonomous operation claims	Trust / regulatory	Medium	High	Use precise language: AI-assisted, pharmacist-overseen, closed pilot; avoid public autonomous dispensing claims	Investor/demo narrative drifts beyond validated scope
Hardware consumes too much schedule	Technical / program	Medium	Medium	Keep hardware to one narrow tablet-dispense or pick-and-place use case; no production robotics claims	RE requires custom fabrication beyond bench rig by week 10
Pilot partner cannot provide enough sample data or pharmacist time	Operational	Medium	High	Secure LOI in week 1; synthetic/blank-form dataset for early build; reserve pharmacist review blocks weekly	Less than 2 hours/week pharmacist time or fewer than 50 representative cases by week 4
ePrescribing integration plan becomes obsolete	Interop / market	Medium	Medium	Build channel-agnostic ingestion and adapter layer; align to Infoway standard, not only PrescribeIT	Vendor or jurisdiction requirements shift during pilot planning

Risk	Category	Likelihood	Impact	Mitigation	Trigger / watchpoint
Patient chatbot gives unsafe open-ended advice	Clinical / trust	Medium	High	Limit to approved content and FAQ; mandatory pharmacist escalation; no diagnosis/prescribing functionality	Unsupported answers appear in UAT
Team tries to include controlled drugs too early	Regulatory / clinical	Low-Medium	High	Hard-scope non-controlled oral solids only for MVP	Pilot partner asks for opioids/controlled workflows in v1

Go or no-go gates

Gate	Timing	Go condition	No-go consequence
Gate A	End of week 4	Accredited partner path, pharmacist availability, representative case set, and security architecture are real	Freeze hardware ambitions; continue as software-only internal prototype
Gate B	End of week 8	End-to-end intake to pharmacist-review workflow works with acceptable accuracy on typed/faxed inputs	Remove counseling and inventory ambitions; focus only on extraction + dashboard
Gate C	End of week 12	Auditability, interaction screening, and narrow hardware demo all function reliably	Defer robotics; ship software-only ops MVP
Gate D	End of week 16	Dry-run metrics, security checks, SOPs, and pharmacist UAT indicate pilot readiness	Do not claim pilot-ready; reframe as internal alpha

Open questions and limitations

The most important unresolved variables are not technical. They are operational:

- **Do you already have an accredited Ontario pharmacy partner or only a concept?** If not, weeks 1–4 must heavily prioritize partner acquisition.
- **Will the supervising pharmacist be onsite in the parent pharmacy during any automated-dispensing pilot?** In Ontario, that matters materially for remote-dispensing design. ⁹
- **Which prescriber classes matter at launch?** Physicians and NPs are enough for a first version, but dental and other prescriber types change registry integration scope. ³⁷
- **Do you need public pilot readiness in 16 weeks, or investor-demo readiness?** Those are different projects. Public pilot readiness has materially more legal, operational, and SOP work.

- **How much real prescription data can be obtained, de-identified, and labeled early?** The accuracy of parsing and authenticity triage will depend more on this than on model choice.
- **Which clinical content vendor can you actually contract inside the first month?** That choice affects schedule risk more than algorithm design.

The highest-confidence conclusion from the available evidence is that a four-month build is viable **only as a tightly-scoped Ontario-compliant internal or partner-site MVP** centered on safe automation, auditable pharmacist oversight, and one narrow dispensing demonstration. Trying to launch a truly “solely AI run pharmacy” in that timeframe would not be a strong regulatory or product strategy in Ontario. ³⁸

¹ ⁸ Opening a Community Pharmacy – OCP website

https://ocpinfo.com/pharmacies/opening-a-community-pharmacy/?utm_source=chatgpt.com

² ⁹ ¹⁰ ¹¹ Remote Dispensing Frequently Asked Questions - OCPInfo.com

https://www.ocpinfo.com/practice-education/opening-operating-pharmacy/opening-rdl/remote-dispensing-faq/?utm_source=chatgpt.com

³ Personal Health Information Protection Act, 2004, S.O. 2004, c. 3, Sched. A"

https://www.ontario.ca/laws/statute/04p03/v47?utm_source=chatgpt.com

⁴ ¹³ ¹⁹ Drug Product Database: Access the database - Canada.ca

https://www.canada.ca/en/health-canada/services/drugs-health-products/drug-products/drug-product-database.html?utm_source=chatgpt.com

⁵ ²⁴ <https://www.prescribit.ca/faq>

<https://www.prescribit.ca/faq>

⁶ ¹⁶ ³⁸ <https://ocpinfo.com/about/about-the-college-and-the-profession/>

<https://ocpinfo.com/about/about-the-college-and-the-profession/>

⁷ We have been receiving emailed prescriptions from prescribers, and sometimes prescribers are emailing prescriptions directly to their patients. Are these emailed prescriptions valid for dispensing? How should we handle these situations? – OCP website

https://ocpinfo.com/ocp-faq/we-have-been-receiving-emailed-prescriptions-from-prescribers-and-sometimes-prescribers-are-emailing-prescriptions-directly-to-their-patients-are-these-emailed-prescriptions-valid-for-dispensing-ho/?utm_source=chatgpt.com

¹² <https://www.aimsclences.org/article/doi/10.3934/jdg.2026016?viewType=HTML>

<https://www.aimsclences.org/article/doi/10.3934/jdg.2026016?viewType=HTML>

¹⁴ <https://ocpinfo.com/pharmacy-professionals/rules-and-standards-of-the-profession/practice-policies-guidelines/operating-internet-sites/>

<https://ocpinfo.com/pharmacy-professionals/rules-and-standards-of-the-profession/practice-policies-guidelines/operating-internet-sites/>

¹⁵ <https://www.ontario.ca/laws/statute/04p03/v38?search=phipa>

<https://www.ontario.ca/laws/statute/04p03/v38?search=phipa>

¹⁷ <https://docs.aws.amazon.com/global-infrastructure/latest/regions/aws-regions.html>

<https://docs.aws.amazon.com/global-infrastructure/latest/regions/aws-regions.html>

¹⁸ <https://arxiv.org/abs/2204.08387>

<https://arxiv.org/abs/2204.08387>

20 CPSO - Public

https://www.cpso.on.ca/en/Public?utm_source=chatgpt.com

21 <https://www.cno.org/protect-the-public/confirm-a-nurses-status>

<https://www.cno.org/protect-the-public/confirm-a-nurses-status>

22 31 https://www.drugbank.com/clinical/drug_drug_interaction_checker

https://www.drugbank.com/clinical/drug_drug_interaction_checker

23 <https://www.rcdso.org/en-ca/standards-guidelines-resources/rcdso-news/articles/14047>

<https://www.rcdso.org/en-ca/standards-guidelines-resources/rcdso-news/articles/14047>

25 30 <https://lhncbc.nlm.nih.gov/RxNav/APIs/index.html>

<https://lhncbc.nlm.nih.gov/RxNav/APIs/index.html>

26 <https://aws.amazon.com/textract/pricing/>

<https://aws.amazon.com/textract/pricing/>

27 <https://cloud.google.com/document-ai/pricing>

<https://cloud.google.com/document-ai/pricing>

28 <https://azure.microsoft.com/en-us/pricing/details/document-intelligence/>

<https://azure.microsoft.com/en-us/pricing/details/document-intelligence/>

29 <https://github.com/tesseract-ocr>

<https://github.com/tesseract-ocr>

32 <https://www.fdbhealth.com/solutions/medknowledge-drug-database/medknowledge-clinical-modules/drug-drug-interaction>

<https://www.fdbhealth.com/solutions/medknowledge-drug-database/medknowledge-clinical-modules/drug-drug-interaction>

33 <https://www.wolterskluwer.com/en-sg/solutions/medi-span>

<https://www.wolterskluwer.com/en-sg/solutions/medi-span>

34 <https://shop.elephantrobotics.com/en-sg/collections/mycobot-280>

<https://shop.elephantrobotics.com/en-sg/collections/mycobot-280>

35 <https://www.robotshop.com/products/ufactory-6-axis-robot-arm-lite-6>

<https://www.robotshop.com/products/ufactory-6-axis-robot-arm-lite-6>

36 <https://www.dobot.us/product/magician-e6-2/>

<https://www.dobot.us/product/magician-e6-2/>

37 Home

https://register.cpso.on.ca/?utm_source=chatgpt.com